


# Agenda

- |   |                                                 |                                    |    |
|---|-------------------------------------------------|------------------------------------|----|
| 1 | Introduction To Machine Learning                | Reinforcement Learning Definitions | 6  |
| 2 | What is Reinforcement Learning?                 | Reinforcement Learning Concepts    | 7  |
| 3 | Reinforcement Learning with an Analogy          | Markov's Decision Process          | 8  |
| 4 | Reinforcement Learning Process                  | Understanding Q-Learning           | 9  |
| 5 | Reinforcement Learning – Counter strike example | Hands-On                           | 10 |

# What Is Machine Learning?

*Machine learning is a subset of artificial intelligence (AI) which provides machines the ability to learn automatically & improve from experience without being explicitly programmed.*



# Types Of Machine Learning



Supervised Learning




Unsupervised Learning



Reinforcement Learning

# What Is Reinforcement Learning?


*Reinforcement learning is a type of Machine Learning where an agent learns to behave in a environment by performing actions and seeing the results*





# Reinforcement Learning With An Analogy

Scenario 1: Baby starts crawling and makes it to the candy




# Reinforcement Learning With An Analogy

Scenario 1: Baby starts crawling and makes it to the candy




# Reinforcement Learning With An Analogy

Scenario 2: Baby starts crawling but falls due to some hurdle in between



# Reinforcement Learning With An Analogy


Scenario 2: Baby starts crawling but falls due to some hurdle in between





# Reinforcement Learning With An Analogy


Scenario 2: Baby starts crawling but falls due to some hurdle in between



# Reinforcement Learning Process

Reinforcement Learning system is comprised of two main components:

- Agent
- Environment



# Reinforcement Learning Process

---

Reinforcement Learning system is comprised of two main components:

- Agent
- Environment



# Reinforcement Learning Process

---

Reinforcement Learning system is comprised of two main components:

- Agent
- Environment



# Reinforcement Learning Process

Reinforcement Learning system is comprised of two main components:

- Agent
- Environment





# Reinforcement Learning Process

Reinforcement Learning system is comprised of two main components:


- Agent
- Environment



# Reinforcement Learning Process

Reinforcement Learning system is comprised of two main components:

- Agent
- Environment



# Counter Strike Example



1. The RL Agent (Player1) collects state  $S^0$  from the environment
2. Based on the state  $S^0$ , the RL agent takes an action  $A^0$ , initially the action is random
3. The environment is now in a new state  $S^1$
4. RL agent now gets a reward  $R^1$  from the environment
5. The RL loop goes on until the RL agent is dead or reaches the destination

# Reinforcement Learning Definitions



Agent: The RL algorithm that learns from trial and error

Environment: The world through which the agent moves



Action (A): All the possible steps that the agent can take

State (S): Current condition returned by the environment



# Reinforcement Learning Definitions



Reward (R): An instant return from the environment to appraise the last action



Policy ( $\pi$ ): The approach that the agent uses to determine the next action based on the current state



Value (V): The expected long-term return with discount, as opposed to the short-term reward R




Action-value (Q): This similar to Value, except, it takes an extra parameter, the current action (A)



# Reward Maximization


Reward maximization theory states that, *a RL agent must be trained in such a way that, he takes the best action so that the reward is maximum.*



# Exploration & Exploitation

*Exploitation* is about using the already known exploited information to heighten the rewards

*Exploration* is about exploring and capturing more information about an environment




# Markov Decision Process

The mathematical approach for mapping a solution in reinforcement learning is called *Markov Decision Process* (MDP)

The following parameters are used to attain a solution:

- Set of actions,  $A$
- Set of states,  $S$
- Reward,  $R$
- Policy,  $\pi$
- Value,  $V$




# Markov Decision Process – Shortest Path Problem

Goal: Find the shortest path between A and D with minimum possible cost


In this problem,

- Set of states are denoted by nodes i.e. {A, B, C, D}
- Action is to traverse from one node to another {A -> B, C -> D}
- Reward is the cost represented by each edge
- Policy is the path taken to reach the destination {A -> C -> D}



# Understanding Q-Learning With An Example

*Place an agent in any one of the rooms (0,1,2,3,4) and the goal is to reach outside the building (room 5)*





- 5 rooms in a building connected by doors
- each room is numbered 0 through 4
- The outside of the building can be thought of as one big room (5)
- Doors 1 and 4 lead into the building from room 5 (outside)



# Understanding Q-Learning With An Example


Let's represent the rooms on a graph, each room as a node, and each door as a link



# Understanding Q-Learning With An Example

Next step is to associate a reward value to each door:


- doors that lead directly to the goal have a reward of 100
- Doors not directly connected to the target room have zero reward
- Because doors are two-way, two arrows are assigned to each room
- Each arrow contains an instant reward value



# Understanding Q-Learning With An Example

The terminology in Q-Learning includes the terms state and action:

- Room (including room 5) represents a state
- agent's movement from one room to another represents an action
- In the figure, a state is depicted as a node, while "action" is represented by the arrows




Example (Agent traverse from room 2 to room5):

1. Initial state = state 2
2. State 2 -> state 3
3. State 3 -> state (2, 1, 4)
4. State 4 -> state 5

# Understanding Q-Learning With An Example

We can put the state diagram and the instant reward values into a reward table, matrix  $R$ .



State

	Action					
	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100

$R =$

The -1's in the table represent null values



# Understanding Q-Learning With An Example

Add another matrix Q, representing the memory of what the agent has learned through experience.

- The rows of matrix Q represent the current state of the agent
- columns represent the possible actions leading to the next state
- Formula to calculate the Q matrix:

$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max} [Q(\text{next state}, \text{all actions})]$$

## Note

The Gamma parameter has a range of 0 to 1 ( $0 \leq \text{Gamma} < 1$ ).

- If Gamma is closer to zero, the agent will tend to consider only immediate rewards.
- If Gamma is closer to one, the agent will consider future rewards with greater weight



# Q – Learning Algorithm

- 1 Set the gamma parameter, and environment rewards in matrix R
- 2 Initialize matrix Q to zero
- 3 Select a random initial state
- 4 Set initial state = current state
- 5 Select one among all possible actions for the current state
- 6 Using this possible action, consider going to the next state
- 7 Get maximum Q value for this next state based on all possible actions
- 8 Compute:  $Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$
- 9 Repeat above steps until current state = goal state

# Q – Learning Example

First step is to set the value of the learning parameter  $\text{Gamma} = 0.8$ , and the initial state as Room 1.

Next, initialize matrix Q as a zero matrix:


- From room 1 you can either go to room 3 or 5, let's select room 5.
- From room 5, calculate maximum Q value for this next state based on all possible actions:

$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$$

$$Q(1,5) = R(1,5) + 0.8 * \text{Max}[Q(5,1), Q(5,4), Q(5,5)] = 100 + 0.8 * 0 = 100$$

	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	0	0	0	0
3	0	0	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	0

		0	1	2	3	4	5
State	Action	0	1	2	3	4	5
0		-1	-1	-1	-1	0	-1
1		-1	-1	-1	0	-1	100
2		-1	-1	-1	0	-1	-1
3		-1	0	0	-1	0	-1
4		0	-1	-1	0	-1	100
5		-1	0	-1	-1	0	100



# Q – Learning Example

For the next episode, we start with a randomly chosen initial state, i.e. state 3

- From room 3 you can either go to room 1,2 or 4, let's select room 1.
- From room 1, calculate maximum Q value for this next state based on all possible actions:


$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$$

$$Q(3,1) = R(3,1) + 0.8 * \text{Max}[Q(1,3), Q(1,5)] = 0 + 0.8 * [0, 100] = 80$$

The matrix Q get's updated

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 80 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

$$R = \begin{matrix} & \begin{matrix} \text{Action} \\ 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} \text{State} \\ 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$



# Q – Learning Example

For the next episode, the next state, 1, now becomes the current state. We repeat the inner loop of the Q learning algorithm because state 1 is not the goal state.


- From room 1 you can either go to room 3 or 5, let's select room 5.
- From room 5, calculate maximum Q value for this next state based on all possible actions:

$$Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$$

$$Q(1,5) = R(1,5) + 0.8 * \text{Max}[Q(5,1), Q(5,4), Q(5,5)] = 100 + 0.8 * 0 = 100$$

The matrix Q remains the same since, Q(1,5) is already fed to the agent

		0	1	2	3	4	5
Q =	0	0	0	0	0	0	0
	1	0	0	0	0	0	100
	2	0	0	0	0	0	0
	3	0	80	0	0	0	0
	4	0	0	0	0	0	0
	5	0	0	0	0	0	0
		Action					
R =	State	0	1	2	3	4	5
	0	-1	-1	-1	-1	0	-1
	1	-1	-1	-1	0	-1	100
	2	-1	-1	-1	0	-1	-1
	3	-1	0	0	-1	0	-1
	4	0	-1	-1	0	-1	100
	5	-1	0	-1	-1	0	100





```
1 import numpy as np
2
3 # R matrix
4 R = np.matrix([[ -1, -1, -1, -1, 0, -1],
5                [-1, -1, -1, 0, -1, 100],
6                [-1, -1, -1, 0, -1, -1],
7                [-1, 0, 0, -1, 0, -1],
8                [-1, 0, 0, -1, -1, 100],
9                [-1, 0, -1, -1, 0, 100]])
10
11 # Q matrix
12 Q = np.matrix(np.zeros([6, 6]))
13
14 # Gamma (learning parameter).
15 gamma = 0.8
16
17 # Initial state. (Usually to be chosen at random)
18 initial_state = 1
19
20
```



```
16
17 # Initial state. (Usually to be chosen at random)
18 initial_state = 1
19
20
21 # This function returns all available actions in the state given as an argument
22 def available_actions(state):
23     current_state_row = R[state_]
24     av_act = np.where(current_state_row >= 0)[1]
25     return av_act
26
27
28 # Get available actions in the current state
29 available_act = available_actions(initial_state)
30
31
32 # This function chooses at random which action to be performed within the range
33 # of all the available actions.
34 def sample_next_action(available_actions_range):
35     next_action = int(np.random.choice(available_act, 1))
```

```
28 # Get available actions in the current state
29 available_act = available_actions(initial_state)
30
31
32 # This function chooses at random which action to be performed within the range
33 # of all the available actions.
34 def sample_next_action(available_actions_range):
35     next_action = int(np.random.choice(available_act, 1))
36     return next_action
37
38
39 # Sample next action to be performed
40 action = sample_next_action(available_act)
41
42
43 # This function updates the Q matrix according to the path selected and the Q
44 # learning algorithm
45 def update(current_state, action, gamma):
46     max_index = np.where(Q[action_] == np.max(Q[action_]))[1]
47
48     sample_next_action()
```

```

37
38
39 # Sample next action to be performed
40 action = sample_next_action(available_act)
41
42
43 # This function updates the Q matrix according to the path selected and the Q
44 # learning algorithm
45 def update(current_state, action, gamma):
46     max_index = np.where(Q[action,] == np.max(Q[action,]))[1]
47
48     if max_index.shape[0] > 1:
49         max_index = int(np.random.choice(max_index, size=1))
50     else:
51         max_index = int(max_index)
52     max_value = Q[action, max_index]
53
54     # Q learning formula
55     Q[current_state, action] = R[current_state, action] + gamma * max_value
56
57     update()

```

Run: demo ×

▶ Run ⚙️ TODO 📄 Terminal 🐍 Python Console

🔍 Event Log

55:11 CRLF UTF-8 4 spaces



```

51         max_index = int(max_index)
52         max_value = Q[action, max_index]
53
54         # Q learning formula
55         Q[current_state, action] = R[current_state, action] + gamma * max_value
56
57
58     # Update Q matrix
59     update(initial_state, action, gamma)
60
61     # -----
62     # Training
63
64     # Train over 10 000 iterations. (Re-iterate the process above).
65     for i in range(10000):
66         current_state = np.random.randint(0, int(Q.shape[0]))
67         available_act = available_actions(current_state)
68         action = sample_next_action(available_act)
69         update(current_state, action, gamma)
70
71     for i in range(10000)

```

```
69     update(current_state, action, gamma)
70
71     # Normalize the "trained" Q matrix
72     print("Trained Q matrix:")
73     print(Q / np.max(Q) * 100)
74
75     # -----
76     # Testing
77
78     # Goal state = 5
79     # Best sequence path starting from 2 -> 2, 3, 1, 5
80
81     current_state = 1
82     steps = [current_state]
83
84     while current_state != 5:
85
86         next_step_index = np.where(Q[current_state, :] == np.max(Q[current_state, :]))[1]
87
88         if next_step_index.shape[0] > 1:
```



```
81 current_state = 1
82 steps = [current_state]
83
84 while current_state != 5:
85
86     next_step_index = np.where(Q[current_state_] == np.max(Q[current_state_]))[1]
87
88     if next_step_index.shape[0] > 1:
89         next_step_index = int(np.random.choice(next_step_index, size=1))
90     else:
91         next_step_index = int(next_step_index)
92
93     steps.append(next_step_index)
94     current_state = next_step_index
95
96 # Print selected sequence of steps
97 print("Selected path:")
98 print(steps)
```

```

79 # Best sequence path starting from 2 -> 2, 3, 1, 5
80
81 current_state = 1
82 steps = [current_state]
83
84 while current_state != 5:
85

```

Run: demo x

C:\Users\zulaikha\PycharmProjects\Test\venv\Scripts\python.exe C:/Users/zulaikha/PycharmProjects/Test/demo

Trained Q matrix:

[	0.	0.	0.	0.	80.	0.]
[	0.	0.	0.	64.	0.	100.]
[	0.	0.	0.	64.	0.	0.]
[	0.	80.	51.2	0.	80.	0.]
[	0.	80.	51.2	0.	0.	100.]
[	0.	80.	0.	0.	80.	100.]]

Selected path:  
[1, 5]

Process finished with exit code 0

```

79 # Best sequence path starting from 2 -> 2, 3, 1, 5
80
81 current_state = 2
82 steps = [current_state]
83
84 while current_state != 5:
85
86     next_step_index = np.where(Q[current_state,:] == np.max(Q[current_state,:]))[1]
87
88     if next_step_index.shape[0] > 1:
89         next_step_index = int(np.random.choice(next_step_index, size=1))
90     else:
91         next_step_index = int(next_step_index)
92
93     steps.append(next_step_index)
94     current_state = next_step_index
95
96 # Print selected sequence of steps

```

while current\_state != 5

Run: demo x

0.	80.	51.2	0.	80.	0.
0.	80.	51.2	0.	0.	100.
0.	80.	0.	0.	80.	100.

Run TODO Terminal Python Console

Event Log

86:34 CRLF UTF-8 4 spaces



```
79 # Best sequence path starting from 2 -> 2, 3, 1, 5
```

```
80
```

```
81 current_state = 2
```

```
82 steps = [current_state]
```

```
83
```

```
84 while current_state != 5:
```

```
85
```

```
86     next_step_index = np.where(Q[current_state,] == np.max(Q[current_state,]))[1]
```

```
87
```

```
88     if next_step_index.shape[0] > 1:
```

```
89         next_step_index = int(np.random.choice(next_step_index, size=1))
```

```
while current_state != 5
```

Trained Q matrix:

```
[[ 0.  0.  0.  0.  80.  0.]
 [ 0.  0.  0.  64.  0. 100.]
 [ 0.  0.  0.  64.  0.  0.]
 [ 0.  80.  51.2 0.  80.  0.]
 [ 0.  80.  51.2 0.  0. 100.]
 [ 0.  80.  0.  0.  80. 100.]]
```

Selected path:

```
[2, 3, 4, 5]
```